

# COSC 6373 - HW2-ICA - Minh Nguyen

## #2069407

```
In [77]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tensorflow import keras
from keras.applications import ResNet50
from keras.applications.resnet50 import preprocess_input
from keras import layers, models
from sklearn.metrics import confusion_matrix, classification_report
```

```
In [78]: # Define dataset paths
train_dir = "train"
test_dir = "test"

# Define parameters
IMG_SIZE = (256, 256)
BATCH_SIZE = 32

# Load training and validation data
train_data = tf.keras.utils.image_dataset_from_directory(
    train_dir,
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    validation_split=0.1,
    subset="training",
    seed=42
)

val_data = tf.keras.utils.image_dataset_from_directory(
    train_dir,
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    validation_split=0.1,
    subset="validation",
    seed=42
)

# Load test dataset
test_data = tf.keras.utils.image_dataset_from_directory(
    test_dir,
    image_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    shuffle=False
)
```

Found 360 files belonging to 2 classes.  
Using 324 files for training.  
Found 360 files belonging to 2 classes.  
Using 36 files for validation.  
Found 40 files belonging to 2 classes.

```
In [79]: # Apply ResNet50 preprocessing
train_data = train_data.map(lambda x, y: (preprocess_input(x), y))
val_data = val_data.map(lambda x, y: (preprocess_input(x), y))

# Load ResNet50 as the base model
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(256, 256, 3))
base_model.trainable = False # Freeze base model initially
```

```
In [80]: # Add custom layers on top
model = models.Sequential([
    base_model,
    layers.GlobalAveragePooling2D(),
    layers.Dropout(0.5), # Prevent overfitting
    layers.Dense(1, activation='sigmoid') # Binary classification
])
```

```
In [81]: # Compile model
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),
              loss='binary_crossentropy',
              metrics=['accuracy'])

# Train initial model
history = model.fit(train_data, validation_data=val_data, epochs=5)
```

```
In [82]: # Unfreeze some layers for fine-tuning
base_model.trainable = True

# Unfreeze the last 25 layers to fine-tune the model
for layer in base_model.layers[-25:]:
    layer.trainable = True

# Recompile model with lower learning rate
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5),
              loss='binary_crossentropy',
              metrics=['accuracy'])

# Model summary
model.summary()

# Fine-tune model
history_fine = model.fit(train_data, validation_data=val_data, epochs=15)
```









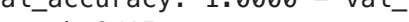
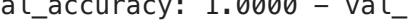
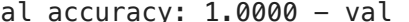
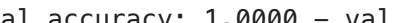
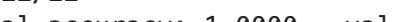


**Model: "sequential\_11"**

Layer (type)	Output Shape	Par
resnet50 (Functional)	(None, 8, 8, 2048)	23,587
global_average_pooling2d_11 (GlobalAveragePooling2D)	(None, 2048)	
dropout_1 (Dropout)	(None, 2048)	
dense_11 (Dense)	(None, 1)	2

**Total params:** 23,589,761 (89.99 MB)

**Trainable params:** 9,992,193 (38.12 MB)

**Non-trainable params:** 13,597,568 (51.87 MB)

Epoch 1/15  
**11/11**  **41s** 3s/step - accuracy: 0.4814 - loss: 0.8856 - val\_accuracy: 0.6944 - val\_loss: 0.6365  
Epoch 2/15  
**11/11**  **30s** 3s/step - accuracy: 0.6177 - loss: 0.6515 - val\_accuracy: 0.7500 - val\_loss: 0.5650  
Epoch 3/15  
**11/11**  **33s** 3s/step - accuracy: 0.7050 - loss: 0.5537 - val\_accuracy: 0.7778 - val\_loss: 0.5008  
Epoch 4/15  
**11/11**  **31s** 3s/step - accuracy: 0.7735 - loss: 0.4848 - val\_accuracy: 0.8333 - val\_loss: 0.4409  
Epoch 5/15  
**11/11**  **30s** 3s/step - accuracy: 0.7971 - loss: 0.4448 - val\_accuracy: 0.8333 - val\_loss: 0.3839  
Epoch 6/15  
**11/11**  **30s** 3s/step - accuracy: 0.8708 - loss: 0.3715 - val\_accuracy: 0.8611 - val\_loss: 0.3364  
Epoch 7/15  
**11/11**  **29s** 3s/step - accuracy: 0.8972 - loss: 0.3083 - val\_accuracy: 0.9444 - val\_loss: 0.2897  
Epoch 8/15  
**11/11**  **29s** 3s/step - accuracy: 0.8936 - loss: 0.3145 - val\_accuracy: 1.0000 - val\_loss: 0.2528  
Epoch 9/15  
**11/11**  **30s** 3s/step - accuracy: 0.9283 - loss: 0.2727 - val\_accuracy: 1.0000 - val\_loss: 0.2190  
Epoch 10/15  
**11/11**  **30s** 3s/step - accuracy: 0.9460 - loss: 0.2056 - val\_accuracy: 1.0000 - val\_loss: 0.1859  
Epoch 11/15  
**11/11**  **30s** 3s/step - accuracy: 0.9680 - loss: 0.1910 - val\_accuracy: 1.0000 - val\_loss: 0.1576  
Epoch 12/15  
**11/11**  **28s** 3s/step - accuracy: 0.9782 - loss: 0.1586 - val\_accuracy: 1.0000 - val\_loss: 0.1372  
Epoch 13/15  
**11/11**  **16s** 1s/step - accuracy: 0.9674 - loss: 0.1441 - val\_accuracy: 1.0000 - val\_loss: 0.1197  
Epoch 14/15  
**11/11**  **15s** 1s/step - accuracy: 0.9954 - loss: 0.1036 - val\_accuracy: 1.0000 - val\_loss: 0.1088  
Epoch 15/15  
**11/11**  **15s** 1s/step - accuracy: 0.9883 - loss: 0.1026 - val\_accuracy: 1.0000 - val\_loss: 0.0968


```
In [83]: # Evaluate model on test data
test_loss, test_accuracy = model.evaluate(test_data)
print(f"Test Accuracy: {test_accuracy:.4f}")

# Generate confusion matrix
y_pred = (model.predict(test_data) > 0.5).astype(int)
y_true = np.concatenate([y for x, y in test_data], axis=0)
```

2/2 ————— 1s 283ms/step – accuracy: 0.9500 – loss: 0.2286  
Test Accuracy: 0.9250  
2/2 ————— 3s 1s/step

```
In [84]: conf_matrix = confusion_matrix(y_true, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Ca
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# Classification report
print("\nClassification Report:\n", classification_report(y_true, y_pred, ta
```



Classification Report:				
	precision	recall	f1-score	support
Camel	0.87	1.00	0.93	20
Horse	1.00	0.85	0.92	20
accuracy			0.93	40
macro avg	0.93	0.93	0.92	40
weighted avg	0.93	0.93	0.92	40

Summary:

- After several attempts with different number of unfreezed layers and number of epochs, I ended up with the following configuration:
  - Unfreezed the last 25 layers
  - Epochs: 15
  - Learning rate:  $1e-5$
- It obtained an accuracy of 0.9500 and a loss of 0.2286 on the test dataset.
- The confusion matrix shows that the model performs pretty well on the test dataset, with only 3 misclassified images out of 40 (3 camels were misclassified as horses).