

# COSC 6373 - HW3-ICA - Minh Nguyen

## #2069407

```
In [1]: import numpy as np
import keras
from keras import layers
```

```
In [4]: # Model / data parameters
num_classes = 10
input_shape = (28, 28, 1)
```

```
In [8]: # Load the data and split it between train and test sets
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# split 10% from the training data for validation
split_index = int(len(x_train) * 0.9)
x_val, x_train = x_train[split_index:], x_train[:split_index]
y_val, y_train = y_train[split_index:], y_train[:split_index]

# Scale images to the [0, 1] range
x_train = x_train.astype("float32") / 255
x_val = x_val.astype("float32") / 255
x_test = x_test.astype("float32") / 255

# Make sure images have shape (28, 28, 1)
x_train = np.expand_dims(x_train, -1)
x_val = np.expand_dims(x_val, -1)
x_test = np.expand_dims(x_test, -1)
print("x_train shape:", x_train.shape)
print("x_val shape:", x_val.shape)
print("x_test shape:", x_test.shape)
print(x_train.shape[0], "train samples")
print(x_val.shape[0], "validation samples")
print(x_test.shape[0], "test samples")

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_val = keras.utils.to_categorical(y_val, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
x_train shape: (54000, 28, 28, 1)
x_val shape: (6000, 28, 28, 1)
x_test shape: (10000, 28, 28, 1)
54000 train samples
6000 validation samples
10000 test samples
```

```
In [6]: # Define the model
model = keras.Sequential(
    [
```

```

keras.Input(shape=input_shape),
layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
layers.MaxPooling2D(pool_size=(2, 2)),
layers.Flatten(),
layers.Dropout(0.5),
layers.Dense(num_classes, activation="softmax"),
]
)

model.summary()

```

**Model: "sequential\_1"**

Layer (type)	Output Shape	Par
conv2d_2 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 32)	
flatten_2 (Flatten)	(None, 5408)	
dropout_2 (Dropout)	(None, 5408)	
dense_1 (Dense)	(None, 10)	54

**Total params:** 54,410 (212.54 KB)

**Trainable params:** 54,410 (212.54 KB)

**Non-trainable params:** 0 (0.00 B)

```

In [9]: # Train the model for 10 epochs (batch_size=32, optimizer=Adam, dropout=0.5,
batch_size = 32
epochs = 10

model.compile(loss="categorical_crossentropy", optimizer="adam", metrics=["a
model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, validation

```

Epoch 1/10  
**1688/1688** ————— **8s** 5ms/step – accuracy: 0.8597 – loss: 0.4938  
 – val\_accuracy: 0.9717 – val\_loss: 0.1075  
 Epoch 2/10  
**1688/1688** ————— **9s** 5ms/step – accuracy: 0.9577 – loss: 0.1425  
 – val\_accuracy: 0.9782 – val\_loss: 0.0808  
 Epoch 3/10  
**1688/1688** ————— **10s** 6ms/step – accuracy: 0.9649 – loss: 0.116  
 1 – val\_accuracy: 0.9820 – val\_loss: 0.0697  
 Epoch 4/10  
**1688/1688** ————— **9s** 6ms/step – accuracy: 0.9728 – loss: 0.0928  
 – val\_accuracy: 0.9828 – val\_loss: 0.0600  
 Epoch 5/10  
**1688/1688** ————— **10s** 6ms/step – accuracy: 0.9754 – loss: 0.082  
 3 – val\_accuracy: 0.9845 – val\_loss: 0.0580  
 Epoch 6/10  
**1688/1688** ————— **10s** 6ms/step – accuracy: 0.9773 – loss: 0.074  
 4 – val\_accuracy: 0.9865 – val\_loss: 0.0520  
 Epoch 7/10  
**1688/1688** ————— **11s** 7ms/step – accuracy: 0.9793 – loss: 0.066  
 4 – val\_accuracy: 0.9872 – val\_loss: 0.0482  
 Epoch 8/10  
**1688/1688** ————— **10s** 6ms/step – accuracy: 0.9811 – loss: 0.059  
 1 – val\_accuracy: 0.9862 – val\_loss: 0.0515  
 Epoch 9/10  
**1688/1688** ————— **10s** 6ms/step – accuracy: 0.9809 – loss: 0.058  
 6 – val\_accuracy: 0.9865 – val\_loss: 0.0502  
 Epoch 10/10  
**1688/1688** ————— **9s** 5ms/step – accuracy: 0.9822 – loss: 0.0545  
 – val\_accuracy: 0.9877 – val\_loss: 0.0469

Out[9]: <keras.src.callbacks.history.History at 0x1674ac2f0>


```
In [10]: # Evaluate the model on the test data
score = model.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score[0])
print("Test accuracy:", score[1])
```

Test loss: 0.04514104872941971  
 Test accuracy: 0.9857000112533569

```
In [11]: # Report the confusion matrix and the test accuracy
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

y_pred = model.predict(x_test)
y_pred_class = np.argmax(y_pred, axis=1)
y_true_class = np.argmax(y_test, axis=1)
confusion_mtx = confusion_matrix(y_true_class, y_pred_class)
sns.heatmap(confusion_mtx, annot=True, fmt="d")
plt.show()
print("Test accuracy:", np.mean(y_true_class == y_pred_class))
```

**313/313** ————— **1s** 2ms/step



Test accuracy: 0.9857

```
In [12]: # Update the base model by adding more layers - x2 Conv2D layers
model_x2_conv = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model_x2_conv.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Par
conv2d_3 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_3 (MaxPooling2D)	(None, 13, 13, 32)	
conv2d_4 (Conv2D)	(None, 11, 11, 64)	18
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 64)	
flatten_3 (Flatten)	(None, 1600)	
dropout_3 (Dropout)	(None, 1600)	
dense_2 (Dense)	(None, 10)	16

**Total params:** 34,826 (136.04 KB)

**Trainable params:** 34,826 (136.04 KB)

**Non-trainable params:** 0 (0.00 B)

```
In [13]: # Train the updated model for 10 epochs (batch_size=32, optimizer=Adam, drop
model_x2_conv.compile(loss="categorical_crossentropy", optimizer="adam", met
model_x2_conv.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, va
```

Epoch 1/10

1688/1688 ————— 14s 8ms/step - accuracy: 0.8496 - loss: 0.472  
6 - val\_accuracy: 0.9818 - val\_loss: 0.0618

Epoch 2/10

1688/1688 ————— 14s 8ms/step - accuracy: 0.9716 - loss: 0.092  
0 - val\_accuracy: 0.9885 - val\_loss: 0.0399

Epoch 3/10

1688/1688 ————— 13s 8ms/step - accuracy: 0.9782 - loss: 0.069  
2 - val\_accuracy: 0.9907 - val\_loss: 0.0346

Epoch 4/10

1688/1688 ————— 14s 8ms/step - accuracy: 0.9817 - loss: 0.057  
1 - val\_accuracy: 0.9902 - val\_loss: 0.0354

Epoch 5/10

1688/1688 ————— 14s 8ms/step - accuracy: 0.9836 - loss: 0.050  
7 - val\_accuracy: 0.9910 - val\_loss: 0.0331

Epoch 6/10

1688/1688 ————— 13s 8ms/step - accuracy: 0.9850 - loss: 0.046  
3 - val\_accuracy: 0.9918 - val\_loss: 0.0295

Epoch 7/10

1688/1688 ————— 13s 8ms/step - accuracy: 0.9865 - loss: 0.043  
1 - val\_accuracy: 0.9927 - val\_loss: 0.0290

Epoch 8/10

1688/1688 ————— 14s 8ms/step - accuracy: 0.9881 - loss: 0.037  
2 - val\_accuracy: 0.9930 - val\_loss: 0.0278

Epoch 9/10

1688/1688 ————— 14s 8ms/step - accuracy: 0.9878 - loss: 0.038  
2 - val\_accuracy: 0.9925 - val\_loss: 0.0289

Epoch 10/10

1688/1688 ————— 13s 8ms/step - accuracy: 0.9900 - loss: 0.032  
0 - val\_accuracy: 0.9920 - val\_loss: 0.0261

Out[13]: <keras.src.callbacks.history.History at 0x31793c830>


```
In [15]: # Evaluate the updated model on the test data
score_x2_conv = model_x2_conv.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score_x2_conv[0])
print("Test accuracy:", score_x2_conv[1])

# Report the confusion matrix and the test accuracy
y_pred_x2_conv = model_x2_conv.predict(x_test)
y_pred_class_x2_conv = np.argmax(y_pred_x2_conv, axis=1)
confusion_mtx_x2_conv = confusion_matrix(y_true_class, y_pred_class_x2_conv)
sns.heatmap(confusion_mtx_x2_conv, annot=True, fmt="d")
plt.show()
print("Test accuracy:", np.mean(y_true_class == y_pred_class_x2_conv))
```

Test loss: 0.025922689586877823

Test accuracy: 0.9911999702453613

313/313 ————— 1s 3ms/step



Test accuracy: 0.9912

- The updated model with 2 Conv2D layers performs slightly better than the base model. It has a higher accuracy and lower loss. (Accuracy: 0.9912 vs 0.9867, Loss: 0.02 vs 0.04)

```
In [16]: # Update the base model by adding more layers - x3 Conv2D layers
model_x3_conv = keras.Sequential([
    keras.Input(shape=input_shape),
    layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
```

```

        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Conv2D(128, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model_x3_conv.summary()

```

**Model: "sequential\_3"**

Layer (type)	Output Shape	Par
conv2d_5 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_5 (MaxPooling2D)	(None, 13, 13, 32)	
conv2d_6 (Conv2D)	(None, 11, 11, 64)	18
max_pooling2d_6 (MaxPooling2D)	(None, 5, 5, 64)	
conv2d_7 (Conv2D)	(None, 3, 3, 128)	73
max_pooling2d_7 (MaxPooling2D)	(None, 1, 1, 128)	
flatten_4 (Flatten)	(None, 128)	
dropout_4 (Dropout)	(None, 128)	
dense_3 (Dense)	(None, 10)	1

**Total params:** 93,962 (367.04 KB)








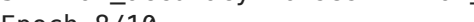
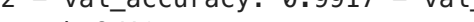

**Trainable params:** 93,962 (367.04 KB)

**Non-trainable params:** 0 (0.00 B)

```

In [17]: # Train the updated model for 10 epochs (batch_size=32, optimizer=Adam, drop
model_x3_conv.compile(loss="categorical_crossentropy", optimizer="adam", met
model_x3_conv.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, va


```

Epoch 1/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.8061 - loss: 0.6048 - val\_accuracy: 0.9808 - val\_loss: 0.0671  
Epoch 2/10  
**1688/1688**  **14s** 9ms/step - accuracy: 0.9618 - loss: 0.1296 - val\_accuracy: 0.9843 - val\_loss: 0.0527  
Epoch 3/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9718 - loss: 0.0932 - val\_accuracy: 0.9873 - val\_loss: 0.0419  
Epoch 4/10  
**1688/1688**  **14s** 9ms/step - accuracy: 0.9777 - loss: 0.0764 - val\_accuracy: 0.9883 - val\_loss: 0.0403  
Epoch 5/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9816 - loss: 0.0610 - val\_accuracy: 0.9888 - val\_loss: 0.0388  
Epoch 6/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9822 - loss: 0.0584 - val\_accuracy: 0.9892 - val\_loss: 0.0379  
Epoch 7/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9844 - loss: 0.0513 - val\_accuracy: 0.9897 - val\_loss: 0.0401  
Epoch 8/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9870 - loss: 0.0422 - val\_accuracy: 0.9917 - val\_loss: 0.0379  
Epoch 9/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9880 - loss: 0.0379 - val\_accuracy: 0.9903 - val\_loss: 0.0399  
Epoch 10/10  
**1688/1688**  **15s** 9ms/step - accuracy: 0.9874 - loss: 0.0398 - val\_accuracy: 0.9888 - val\_loss: 0.0412


Out[17]: <keras.src.callbacks.history.History at 0x317c0c5c0>

```
In [18]: # Evaluate the updated model on the test data
score_x3_conv = model_x3_conv.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score_x3_conv[0])
print("Test accuracy:", score_x3_conv[1])

# Report the confusion matrix and the test accuracy
y_pred_x3_conv = model_x3_conv.predict(x_test)
y_pred_class_x3_conv = np.argmax(y_pred_x3_conv, axis=1)
confusion_mtx_x3_conv = confusion_matrix(y_true_class, y_pred_class_x3_conv)
sns.heatmap(confusion_mtx_x3_conv, annot=True, fmt="d")
plt.show()
print("Test accuracy:", np.mean(y_true_class == y_pred_class_x3_conv))
```

Test loss: 0.038025468587875366  
Test accuracy: 0.9900000095367432  
**313/313**  **1s** 4ms/step





Test accuracy: 0.99

- The model with 3 layers of Conv2D performs similarly to the model with 2 layers. The test accuracy is kind of the same, however, it has a higher loss (0.038 vs 0.02).

```
In [22]: # Update the base model with a different activation function - sigmoid
model_sigmoid = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="sigmoid"),
    ]
)

model_sigmoid.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Par
conv2d_9 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_9 (MaxPooling2D)	(None, 13, 13, 32)	
flatten_6 (Flatten)	(None, 5408)	
dropout_6 (Dropout)	(None, 5408)	
dense_5 (Dense)	(None, 10)	54


**Total params:** 54,410 (212.54 KB)

**Trainable params:** 54,410 (212.54 KB)


**Non-trainable params:** 0 (0.00 B)

```
In [23]: # Train the updated model for 10 epochs (batch_size=32, optimizer=Adam, drop
model_sigmoid.compile(loss="categorical_crossentropy", optimizer="adam", met
model_sigmoid.fit(x_train, y_train, batch_size=batch_size, epochs=epochs, va
```


Epoch 1/10

1688/1688  10s 6ms/step - accuracy: 0.8544 - loss: 0.499  
0 - val\_accuracy: 0.9740 - val\_loss: 0.1008


Epoch 2/10

1688/1688  9s 6ms/step - accuracy: 0.9599 - loss: 0.1361  
- val\_accuracy: 0.9797 - val\_loss: 0.0738


Epoch 3/10

1688/1688  9s 6ms/step - accuracy: 0.9680 - loss: 0.1045  
- val\_accuracy: 0.9828 - val\_loss: 0.0651


Epoch 4/10

1688/1688  9s 6ms/step - accuracy: 0.9724 - loss: 0.0892  
- val\_accuracy: 0.9843 - val\_loss: 0.0605


Epoch 5/10

1688/1688  10s 6ms/step - accuracy: 0.9747 - loss: 0.080  
7 - val\_accuracy: 0.9847 - val\_loss: 0.0589


Epoch 6/10

1688/1688  9s 6ms/step - accuracy: 0.9776 - loss: 0.0719  
- val\_accuracy: 0.9847 - val\_loss: 0.0530


Epoch 7/10

1688/1688  9s 5ms/step - accuracy: 0.9803 - loss: 0.0632  
- val\_accuracy: 0.9850 - val\_loss: 0.0522


Epoch 8/10

1688/1688  10s 6ms/step - accuracy: 0.9810 - loss: 0.061  
2 - val\_accuracy: 0.9848 - val\_loss: 0.0512

Epoch 9/10

1688/1688  9s 6ms/step - accuracy: 0.9824 - loss: 0.0550  
- val\_accuracy: 0.9873 - val\_loss: 0.0492

Epoch 10/10

1688/1688  9s 5ms/step - accuracy: 0.9845 - loss: 0.0496  
- val\_accuracy: 0.9873 - val\_loss: 0.0493

```
Out[23]: <keras.src.callbacks.history.History at 0x317e0a000>
```


```
In [24]: # Evaluate the updated model on the test data
score_sigmoid = model_sigmoid.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score_sigmoid[0])
print("Test accuracy:", score_sigmoid[1])

# Report the confusion matrix and the test accuracy
y_pred_sigmoid = model_sigmoid.predict(x_test)
y_pred_class_sigmoid = np.argmax(y_pred_sigmoid, axis=1)
confusion_mtx_sigmoid = confusion_matrix(y_true_class, y_pred_class_sigmoid)
sns.heatmap(confusion_mtx_sigmoid, annot=True, fmt="d")
plt.show()
print("Test accuracy:", np.mean(y_true_class == y_pred_class_sigmoid))
```

Test loss: 0.0494907945394516

Test accuracy: 0.9843999743461609

313/313 ————— 1s 2ms/step



Test accuracy: 0.9844

- The model with sigmoid activation function performs worse than most of the models above, with a test accuracy of 0.9842 (still very good), but it has a higher loss of 0.05.

```
In [25]: # Update the model and train with a different batch size - 64
model_batch_64 = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
```

```

        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model_batch_64.summary()

```

Model: "sequential\_6"

Layer (type)	Output Shape	Par
conv2d_10 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_10 (MaxPooling2D)	(None, 13, 13, 32)	
flatten_7 (Flatten)	(None, 5408)	
dropout_7 (Dropout)	(None, 5408)	
dense_6 (Dense)	(None, 10)	54

Total params: 54,410 (212.54 KB)

Trainable params: 54,410 (212.54 KB)

Non-trainable params: 0 (0.00 B)

```

In [26]: # Train the updated model for 10 epochs (batch_size=64, optimizer=Adam, drop
batch_size_64 = 64

model_batch_64.compile(loss="categorical_crossentropy", optimizer="adam", me
model_batch_64.fit(x_train, y_train, batch_size=batch_size_64, epochs=epochs

```

```

Epoch 1/10
844/844 ————— 6s 7ms/step - accuracy: 0.8301 - loss: 0.6043 -
val_accuracy: 0.9675 - val_loss: 0.1201
Epoch 2/10
844/844 ————— 6s 7ms/step - accuracy: 0.9548 - loss: 0.1587 -
val_accuracy: 0.9765 - val_loss: 0.0868
Epoch 3/10
844/844 ————— 6s 7ms/step - accuracy: 0.9635 - loss: 0.1198 -
val_accuracy: 0.9818 - val_loss: 0.0711
Epoch 4/10
844/844 ————— 6s 7ms/step - accuracy: 0.9702 - loss: 0.1003 -
val_accuracy: 0.9803 - val_loss: 0.0673
Epoch 5/10
844/844 ————— 6s 7ms/step - accuracy: 0.9724 - loss: 0.0887 -
val_accuracy: 0.9827 - val_loss: 0.0628
Epoch 6/10
844/844 ————— 6s 7ms/step - accuracy: 0.9739 - loss: 0.0839 -
val_accuracy: 0.9833 - val_loss: 0.0590
Epoch 7/10
844/844 ————— 6s 7ms/step - accuracy: 0.9765 - loss: 0.0784 -
val_accuracy: 0.9857 - val_loss: 0.0542
Epoch 8/10
844/844 ————— 6s 7ms/step - accuracy: 0.9776 - loss: 0.0716 -
val_accuracy: 0.9860 - val_loss: 0.0549
Epoch 9/10
844/844 ————— 6s 7ms/step - accuracy: 0.9778 - loss: 0.0696 -
val_accuracy: 0.9858 - val_loss: 0.0536
Epoch 10/10
844/844 ————— 6s 7ms/step - accuracy: 0.9792 - loss: 0.0658 -
val_accuracy: 0.9865 - val_loss: 0.0516

```

Out[26]: <keras.src.callbacks.history.History at 0x317be9730>

```

In [27]: # Evaluate the updated model on the test data
score_batch_64 = model_batch_64.evaluate(x_test, y_test, verbose=0)
print("Test loss:", score_batch_64[0])
print("Test accuracy:", score_batch_64[1])


# Report the confusion matrix and the test accuracy
y_pred_batch_64 = model_batch_64.predict(x_test)
y_pred_class_batch_64 = np.argmax(y_pred_batch_64, axis=1)
confusion_mtx_batch_64 = confusion_matrix(y_true_class, y_pred_class_batch_64)
sns.heatmap(confusion_mtx_batch_64, annot=True, fmt="d")
plt.show()
print("Test accuracy:", np.mean(y_true_class == y_pred_class_batch_64))

```

```

Test loss: 0.051112327724695206
Test accuracy: 0.9835000038146973
313/313 ————— 1s 2ms/step

```



Test accuracy: 0.9835

- The base model but trained with a batch size of 64 has a slightly lower accuracy and higher loss compared to the other models. (Accuracy: 0.9835, Loss: 0.05)

```
In [28]: # Try to train a model with 2 layers of Conv2D with a batch size of 128
model_2x_conv_batch_128 = keras.Sequential(
    [
        keras.Input(shape=input_shape),
        layers.Conv2D(32, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Conv2D(64, kernel_size=(3, 3), activation="relu"),
        layers.MaxPooling2D(pool_size=(2, 2)),
        layers.Flatten(),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

model_2x_conv_batch_128.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Par
conv2d_11 (Conv2D)	(None, 26, 26, 32)	
max_pooling2d_11 (MaxPooling2D)	(None, 13, 13, 32)	
conv2d_12 (Conv2D)	(None, 11, 11, 64)	18
max_pooling2d_12 (MaxPooling2D)	(None, 5, 5, 64)	
flatten_8 (Flatten)	(None, 1600)	
dropout_8 (Dropout)	(None, 1600)	
dense_7 (Dense)	(None, 10)	16

**Total params:** 34,826 (136.04 KB)

**Trainable params:** 34,826 (136.04 KB)

**Non-trainable params:** 0 (0.00 B)

```
In [30]: # Train the updated model for 10 epochs (batch_size=128, optimizer=Adam, dr
model_2x_conv_batch_128.compile(loss="categorical_crossentropy", optimizer="
model_2x_conv_batch_128.fit(x_train, y_train, batch_size=128, epochs=epochs,
```

Epoch 1/10

422/422 ————— 9s 20ms/step - accuracy: 0.9814 - loss: 0.0607  
- val\_accuracy: 0.9907 - val\_loss: 0.0385

Epoch 2/10

422/422 ————— 8s 20ms/step - accuracy: 0.9835 - loss: 0.0521  
- val\_accuracy: 0.9907 - val\_loss: 0.0328

Epoch 3/10

422/422 ————— 8s 18ms/step - accuracy: 0.9844 - loss: 0.0479  
- val\_accuracy: 0.9913 - val\_loss: 0.0316

Epoch 4/10

422/422 ————— 7s 17ms/step - accuracy: 0.9857 - loss: 0.0442  
- val\_accuracy: 0.9918 - val\_loss: 0.0316

Epoch 5/10

422/422 ————— 7s 17ms/step - accuracy: 0.9856 - loss: 0.0460  
- val\_accuracy: 0.9918 - val\_loss: 0.0292

Epoch 6/10

422/422 ————— 8s 20ms/step - accuracy: 0.9874 - loss: 0.0393  
- val\_accuracy: 0.9915 - val\_loss: 0.0301

Epoch 7/10

422/422 ————— 9s 21ms/step - accuracy: 0.9881 - loss: 0.0379  
- val\_accuracy: 0.9917 - val\_loss: 0.0310

Epoch 8/10

422/422 ————— 8s 19ms/step - accuracy: 0.9875 - loss: 0.0375  
- val\_accuracy: 0.9913 - val\_loss: 0.0283

Epoch 9/10

422/422 ————— 8s 19ms/step - accuracy: 0.9885 - loss: 0.0379  
- val\_accuracy: 0.9930 - val\_loss: 0.0273

Epoch 10/10

422/422 ————— 7s 17ms/step - accuracy: 0.9890 - loss: 0.0330  
- val\_accuracy: 0.9933 - val\_loss: 0.0282

Out[30]: <keras.src.callbacks.history.History at 0x343c51130>


```
In [31]: # Evaluate the updated model on the test data
score_2x_conv_batch_128 = model_2x_conv_batch_128.evaluate(x_test, y_test, v
print("Test loss:", score_2x_conv_batch_128[0])
print("Test accuracy:", score_2x_conv_batch_128[1])

# Report the confusion matrix and the test accuracy
y_pred_2x_conv_batch_128 = model_2x_conv_batch_128.predict(x_test)
y_pred_class_2x_conv_batch_128 = np.argmax(y_pred_2x_conv_batch_128, axis=1)
confusion_mtx_2x_conv_batch_128 = confusion_matrix(y_true_class, y_pred_class_2x_conv_batch_128)
sns.heatmap(confusion_mtx_2x_conv_batch_128, annot=True, fmt="d")
plt.show()
print("Test accuracy:", np.mean(y_true_class == y_pred_class_2x_conv_batch_128))
```

Test loss: 0.023889461532235146

Test accuracy: 0.9914000034332275

313/313 ————— 1s 3ms/step



Test accuracy: 0.9914

- The model with 2 layers of Conv2D and a batch size of 128 performs the best among the others in this analysis, with a test accuracy of 0.9914 and a loss of 0.0239.

4. What is the difference between epoch and batch? How does batch size affect the training

time and performance?



- Epoch: epoch is the number of times that the entire training dataset is passed through the network.
- Batch: Batch size is the number of training samples (subset) that are fed to the network at once. The number of batches in an epoch is equal to the number of training samples divided by the batch size.
- The batch size affects the training time and performance in the following ways:
  - Larger batch sizes can speed up the training process because the model updates its weights less frequently per epoch. However, it may lead to a lower generalization performance.
  - Smaller batch sizes can slow down the training process because the model updates its weights more frequently. However, it may lead to a better accuracy by capturing more information from the training data. It may be more computationally expensive.

5. What is the purpose of the max pooling layer?

- The purpose of the max pooling layer is to reduce the spatial dimensions of the input volume. It selects the maximum value within a local region. It helps to reduce the number of parameters and computation in the network, which also can help to prevent overfitting. Max pooling also helps to make the network more robust to small variations in the input data by capturing the most important features in the input volume.